

SUMMARY

Aspiring researcher focused on **computer vision** and **generative AI**, with research experience in **computational photography** and **crime detection**, and work experience at **Stripe**. Skilled in **2D/3D computer vision**, **federated learning**, and **generative modeling** through hands-on research and implementations.

EDUCATION

- Birla Institute of Technology and Science** Pilani, India
Bachelor of Engineering in Computer Science; CGPA: 8.75/10.0 Oct. 2022 – Present
 - Core Coursework:** Data Structures and Algorithms, Database Systems, Object Oriented Programming, Operating Systems
 - Relevant Coursework:** Generative AI, Introduction to LLMs, Natural Language Processing, Deep Learning, Linear Algebra, Probability & Statistics
- Deens Academy, ECC Road** Bengaluru, India
Grade XII - Computer Science; SAT: 1550/1600, CBSE Examinations: 97.6% Mar. 2021 – Jul. 2022

EXPERIENCE

- Stripe** Bengaluru, India
Software Developer Intern May. 2025 – Jul. 2025
 - Saved **\$200K** per annum by reducing **700K** manual verifications for merchant addresses using LLMs.
 - Conducted prompt engineering experiments with **1,000** user addresses to reduce the LLM's false rejection rate to **<2%**.
 - Created an abstract framework to reduce time to create new LLM verification systems from **~ 2 weeks** to **~ 1 day**.
 - Utilised:** Ruby, Typescript, Mongo, Splunk.
- CarScan.AI** Pune, India
Computer Vision Engineer Jun. 2024 – Jul. 2024
 - Used transfer learning on Ultralytics YOLOv8 to recognise **32** car parts for Carscan.ai's vehicle damage detection app.
 - Refined Instance Segmentation by annotating **400+** multi-angle vehicle images using the Computer Vision Annotation Tool.
 - Utilised:** Computer Vision Annotation Tool (CVAT), PyTorch, Python, Kaggle, YOLOv8, and Matplotlib.
- Research on Computational Photography using Federated Learning** Pilani, India
Research Assistant under Dr. Pratik Narang, Dept. of Computer Science, BITS Pilani Aug. 2024 – Present
 - Developing a federated learning solution for advanced digital image processing on RAW images while preserving data privacy.
 - Evaluated federated denoising using **7** noise models, based on **20+** studies, across **13** architectures including DBSN, DnCNN.
 - Trained Blind2Unblind architecture using federated clients on custom noise models, achieving **47.28 dB** PSNR, **0.98** SSIM.
 - Utilised:** PyTorch, Python, WandB, Linux, and Matplotlib.
- Research on Crime Detection in Surveillance Videos** Pilani, India
Research Assistant Nov. 2023 – May 2024
 - Detected and classified **10** crime types in UCF-Crime dataset using YOLOv3 and DeepSORT; preprocessed **1000+** videos.
 - Designed, trained, and tuned an LSTM model, achieving an AUC score of **62.5%**, approaching state-of-the-art performance.
 - Utilised:** Python, PyTorch, Pandas, Scikit-learn, Google Colab, Seaborn, Matplotlib, Numpy, OpenCV, and PIL.

PROJECTS

- 3D Semantic Segmentation of EPFL Electron Microscopy Dataset** May. 2025 – Jul. 2025
 - Experimented with 3D segmentation task using Morphological Post-Processing, nnU-Net, and Foundation Models.
 - Pre-trained PCT-Net using Volume Fusion on unannotated volume, achieving Dice of **99.93%** on pretext task.
 - Analysed and experimented with preprocessing, Focal Dice, schedulers, postprocessing, increasing Dice from **42%** to **93%**.
 - Utilised:** Python, PyMIC, PyTorch, Linux.
- Implementation of Conditional PixelVAE, PixelCNN for Handwritten Digit Generation** Feb. 2025
 - Reproduced PixelCNN, conditional PixelCNN, PixelVAE architectures; extended PixelVAE by creating a conditional variant.
 - Achieved strong quantitative performance on binarised MNIST dataset with average NLL of **0.169** nats per pixel.
 - Utilised:** Python, Google Colab, Numpy, and Matplotlib.
- From-Scratch Neural Learning in NumPy via Adam Optimisation** Jan. 2025
 - Built forward/backward passes with analytical gradients and gradient descent to study neural learning without frameworks.
 - Constructed Adam optimiser from scratch, yielding **100%** test accuracy; visualised model behaviour via decision boundaries.
 - Utilised:** Python, Pandas, Scikit-learn, Google Colab, Numpy, Matplotlib, and Seaborn.
- Analysis of Smoking and Drinking Effects on Health Metrics** Sep. 2023
 - Performed rigorous preprocessing on **10,000+** health records and conducted EDA to identify trends across **15+** metrics.
 - Boosted accuracy from **54%** to **72%** by evaluating models like Logistic Regression, Random Forest; gained **33%** via tuning.
 - Utilised:** Python, Pandas, Scikit-learn, Google Colab, Numpy, Matplotlib, and Seaborn.

PROGRAMMING SKILLS

- Languages:** Python, Ruby, Java, C/C++, SQL
- Technologies:** PyTorch, NumPy, Matplotlib, Keras, TensorFlow